

Homework 3 - 4/25/2019

Monday, April 22, 2019 1:10 PM

Due 4/25/2019

- 1) Problem 4.25: Modifying the socket-based date time server to be multithreaded from Figure 3.21. MAKE SURE YOU TURN in executable .java files
Create two versions
 - a) Creating a new thread for each request
 - b) Using a Java-based thread-pool
- 2) Problem 4.17
- 3) Read and summarize following page: <http://www.thegeekstuff.com/2013/11/linux-process-and-threadsLinks to an external site.>

Exercise 4.17

The program shown in Figure 4.16 uses the PthreadsAPI. What would be the output from the program at LINE C and LINE P?

```
#include <pthread.h>
#include <stdio.h>

#include <types.h>

int value = 0;
void *runner(void *param); /* the thread */

int main(int argc, char *argv[])
{
    pid_t pid;
    pthread_t tid;
    pthread_attr_t attr;

    pid = fork();

    if (pid == 0) { /* child process */
        pthread_attr_init(&attr);
        pthread_create(&tid,&attr,runner,NULL);
        pthread_join(tid,NULL);
        printf("CHILD: value = %d",value); /* LINE C */
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("PARENT: value = %d",value); /* LINE P */
    }
}

void *runner(void *param) {
    value = 5;
    pthread_exit(0);
}
```

Figure 4.16 C program for Exercise 4.17.

LINE C: CHILD: 5

LINE P: PARENT: 0

The value of the global variable *value* is changed in the CHILD program to be 5 in the new thread

created from the *pthread create* command and using the **runner* function. Inside the CHILD memory the *value* would be considered 5, however; the PARENT memory global value for *value* wouldn't be changed. The CHILD and PARENT only share memory when the memory written is explicitly made to be shared memory. Otherwise the memories are separate copies.

Exercise 4.25

Modify the socket-based date server (Figure 3.21) in Chapter 3 so that the server services each client request in a separate thread

```
import java.net.*;
import java.io.*;

public class DateServer
{
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            /* now listen for connections */
            while (true) {
                Socket client = sock.accept();

                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                /* write the Date to the socket */
                pout.println(new java.util.Date().toString());

                /* close the socket and resume */
                /* listening for connections */
                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

Figure 3.21 Date server.

Assume client is using DataClient.java as provided.

Code for part a, new thread per request:

DateServer

```
/*
 * Created by rtdimpsey on 4/13/17.
 * Edited by Daniel Yan 4/25/2019
```

```

*/
import java.net.*;
import java.io.*;

//This multi-threaded implementation uses Java Thread class
public class DateServer extends Thread
{
    private Socket client;
    //Each thread id should be corresponding to another client
    private static int id;
    private final static int MAX_THREAD = 100;
    //Quick constructor to create a DateServer
    public DateServer(Socket client)
    {
        super();
        this.client = client;
    }
    public void run()
    {
        //Try catch block is needed for exception
        while(Thread.activeCount() < MAX_THREAD)
        {
            System.out.println("Running thread[" + Thread.activeCount() + "]");
            try
            {
                //Print out client as a Date object
                PrintWriter pout = new PrintWriter(client.getOutputStream(), true);
                pout.println(new java.util.Date().toString());
                Thread.sleep(1000);
            }
            catch(Exception ie)
            {
                System.err.println(ie);
            }
        }
        //Ending thread print
        System.out.println("All threads are running...for regular Threads");
    }
    public static void main(String[] args)
    {
        try
        {
            //Create a new DateServer thread
            DateServer thread[] = new DateServer[MAX_THREAD];
            ServerSocket sock = new ServerSocket(6014);
            //Create an accept socket
            Socket client = sock.accept();
            for (int i = 0; i < MAX_THREAD; i++)
            {
                //Create a new DateServer thread with id
                thread[i] = new DateServer(client);
                //Run the thread
                thread[i].start();
            }
        }
    }
}

```

```

    }
    //Catch any errant exception
    catch (IOException ie) {
        System.err.println(ie);
    }
}
}

```

Code for part b, thread pool usage

DateServerPool

```

/**
 * Created by rtdimpsey on 4/13/17.
 * Edited by Daniel Yan 4/25/2019
 */
import java.net.*;
import java.io.*;
//Needed ThreadPool import
import java.util.concurrent.Executors;
import java.util.concurrent.ExecutorService;

//This thread pool implementation uses the Java Runnable interface
public class DateServerPool implements Runnable
{
    private Socket client;
    //Each thread id should be corresponding to another client
    private final static int MAX_THREAD = 100;
    //Quick constructor to create a DateServerPool
    public DateServerPool(Socket client)
    {
        super();
        this.client = client;
    }
    public void run(){
        //Try catch block is needed for exception
        while(Thread.activeCount() < MAX_THREAD)
        {
            System.out.println("Running thread[" + Thread.activeCount() + "]");
            try
            {
                //Print the client out as a Date object
                PrintWriter pout = new PrintWriter(client.getOutputStream(), true);
                pout.println(new java.util.Date().toString());
            }
            catch(IOException ie)
            {
                System.err.println(ie);
            }
        }
        System.out.println("All threads are running...for Executor ThreadPool");
    }
}

```

```

public static void main(String[] args)
{
    //Array needed for ExecutorService
    Runnable[] dsp = new DateServerPool[MAX_THREAD];
    try
    {
        ServerSocket sock = new ServerSocket(6014);
        //Make a accept socket
        Socket client = sock.accept();
        ExecutorService threadExecutor = Executors.newFixedThreadPool(MAX_THREAD);
        //As we specified the MAX_THREAD size in array we need a for loop
        for(int i = 0; i < MAX_THREAD; i++)
        {
            //Initialize the thread with client
            dsp[i] = new DateServerPool(client);
            //Execute the thread
            threadExecutor.execute(dsp[i]);
        }
    }
    //For any odd IO exception
    catch (IOException ie) {
        System.err.println(ie);
    }
}

```

Date Client

```

/**
 * Created by rtdimpsey on 4/13/17.
 */
import java.net.*;
import java.io.*;

public class DateClient
{
    public static void main(String[] args)
    {
        try
        {
            Socket sock = new Socket("127.0.0.1",6014);
            InputStream in = sock.getInputStream();
            BufferedReader bin = new BufferedReader(new InputStreamReader(in));
            String line;
            while ((line = bin.readLine()) != null)
            {
                System.out.println(line);
            }
            sock.close();
        }
        catch (IOException ie) {}
    }
}

```

GeekStuff Summarization:

Linux Processes:

- Simply a running instance of a program, where a kernel provides system resources for the process
- Processes have priority assigned based on kernel context switches, and can be pre-empted due to higher priority processes
- Interprocess communication is used for processes to talk to another, and shared memory is used for sharing data
- Fork() creates a child and parent process. Where in the child process if a change is made, a separate copy of stack and the heap are made for the child

Linux Threads & Light Weight Processes:

- Simply a flow of execution for a process, where a process with multiple flows would be called multi-threaded process
- Linux kernel views threads as processes inside of a kernel, these are termed Light Weight Processes (LWPs)
- LWPs share the same address space and other resources comparative to normal processes
- Threads and LWPs are the same, the terminology differentiation is threads for user view and LWPs for kernel view

Linux Process States:

- Each Linux Process will contain various states
- RUNNING: Process is executing or waiting to be executed
- INTERRUPTIBLE: Process is in sleep mode and is waiting to be woken up by interruption
- UN-INTERRUPTIBLE: Process is in sleep mode but cannot be woken up by a signal
- STOPPED: Process is stopped, may be due to a signal (SIGSTOP for example)
- TRACED: Process is being debugged
- ZOMBIE: Process is terminated but parent of process hasn't found the termination of their child
- DEAD: Process is terminated and process table entry is removed. Parent has found terminated child state